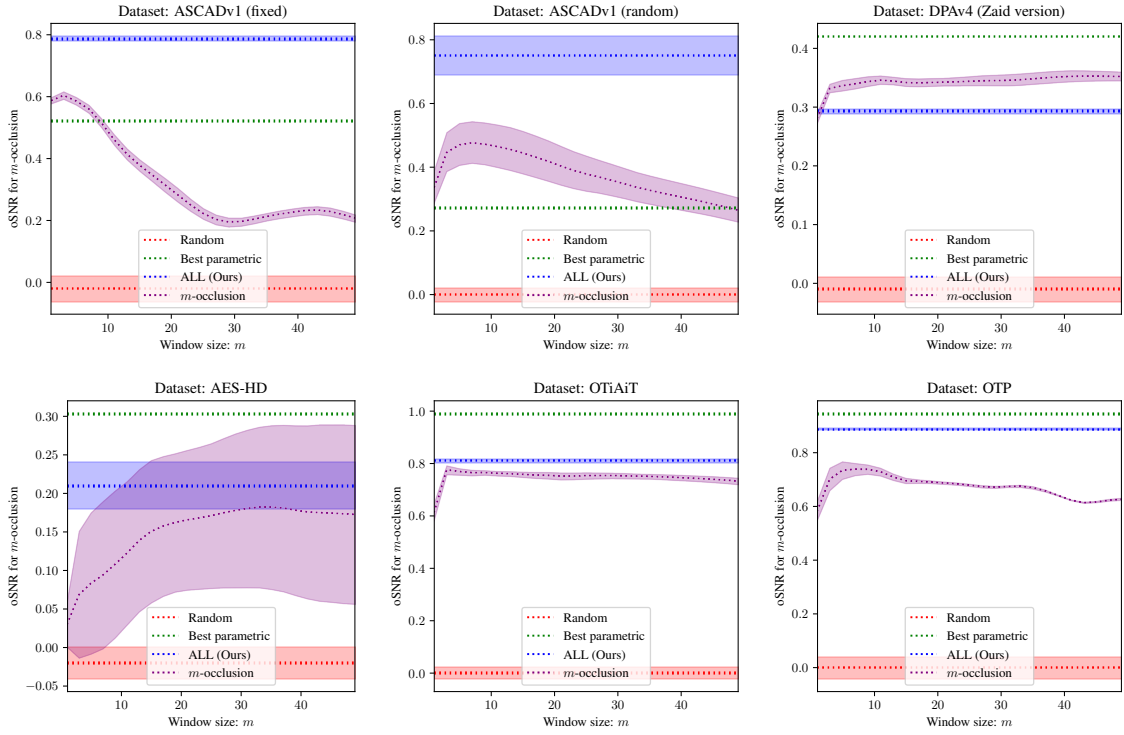

Figure 1: Performance of m-occlusion (purple trace, NOT ours) for occlusion window diameter $m \in \{1, 3, 5, \ldots, 29\}$, measured using the oSNR metric (**larger is better**). Note that the set of valid values for m is $\{1, 2, \ldots, T\}$ where T is the data dimensionality. For reference, we also plot horizontal lines and error bars corresponding to the performance of random guessing, the *best* parametric method, and our ALL algorithm, as indicated in the legend. Observe that while $m > 1$ is optimal for all datasets, for every dataset except AES-HD tuning m provides only a modest performance boost. Performance remains below our ALL algorithm despite tuning except on the DPAv4 dataset. Dotted lines denote mean and shading denotes median over 5 random seeds.